

Object Classification on AWS

Gianmarco Montillo, Simone Boesso, Arianna Celli

1. Introduction

In recent times, the world has witnessed significant challenges and changes in response to the global COVID-19 pandemic. The necessity of wearing masks to prevent the spread of the virus has become a crucial aspect of public health measures.

As a result, the development of efficient and accurate methods for identifying individuals wearing masks has gained paramount importance in various domains, including security, healthcare, and public space management.

AWS provides a wide range of services, including storage (Amazon S3), computing resources (Amazon EC2), and serverless computing (AWS Lambda). These services offer high scalability, flexibility, and reliability, enabling the architecture to effectively handle varying workloads and seamlessly accommodate spikes in demand.

The availability of comprehensive documentation, tutorials, and support from AWS simplifies the implementation and management of the architecture, ensuring a smooth development process. Additionally, AWS ensures a secure and compliant infrastructure, assuring the privacy and security of the data being processed.

Leveraging the AWS approach empowers organizations to build and deploy highly efficient, scalable, and cost-effective object classification solutions in a professional manner.

This report presents a comprehensive overview of our cloud-based architecture designed to address the task of object classification, specifically focused on identifying individuals wearing masks.

This project aims to develop a robust and scalable solution capable of accurately detecting masked individuals in real-time.

Throughout this report, we will delve into the technical aspects, methodologies, and challenges encountered during the implementation of this cloud architecture. Additionally, we will discuss the rationale behind the selected components and technologies, highlighting the advantages they offer for this specific object detection task.

2. Dataset

The dataset ([Face Mask Detection | Kaggle](#)) used for training and testing the mask-wearing behavior classification model consists of 853 images. Each image is in JPEG format and represents individuals wearing or not wearing face masks.

The images are labeled with one of the following three classes:

- *With mask*: Depicts individuals correctly wearing a face mask.
- *Without mask*: Shows individuals not wearing a face mask.
- *Mask worn incorrectly*: Illustrates individuals wearing a face mask, but not in the correct manner.



Annotations:

The dataset provides corresponding XML files for each image, containing bounding box annotations.

These annotations define the coordinates of the bounding boxes around the faces and mask regions within the images.

The bounding box annotations are used to identify the regions of interest (ROIs) for mask detection.

Each bounding box annotation corresponds to a single face in the image and provides information about the position and dimensions of the bounding box. These annotations are crucial for training the model to accurately detect and classify mask-wearing behavior.

In summary, the dataset consists of 853 JPEG images, each associated with an XML file containing bounding box annotations. The annotations provide details about the position and size of the bounding boxes, enabling the model to identify and classify mask-wearing behavior accurately.

3.Application

Our project aimed to develop a pipeline using AWS technologies to predict if a person in an image belongs to one of three classes: wearing a mask, not wearing a mask, or wearing a mask incorrectly.

Initially, we explored two approaches for achieving this task: utilizing a YOLOv5 architecture and a Resnet34 architecture.

However, we encountered numerous challenges while attempting to install the dependencies required for the YOLOv5 architecture.

After successfully setting up our Jupyter notebook, we faced difficulties with installing Torch, and even after resolving this issue, we were unable to proceed with the training phase due to the inability to access the already installed libraries in the environment.

Consequently, we decided to proceed with the Resnet34 architecture and performed training using the following data splits: 60% of the dataset was allocated for the training phase 20% of the dataset was used for validation 20% of the dataset was reserved for testing.

The results obtained by this approach for the training will be enlisted later in the report.

Despite the challenges encountered during the project, we were able to successfully develop the application using the Resnet34 architecture and train it on the provided dataset.

These results mark a significant milestone in achieving our goal of accurately classifying individuals based on their mask-wearing behavior.

4. Infrastructure

Infrastructure and Training Process

Our infrastructure is built upon four key services: EC2, S3, Lambda AWS, and DynamoDB.

1. **EC2:**

We utilized EC2 instances with the appropriate roles (LabRole) to access files in our S3 bucket (*'ourcloudcomputingbucket'*). During the training process, we faced challenges related to the resource requirements of the model.

Specifically, we had to vertically scale our EC2 instances by increasing RAM and storage capacity to ensure successful model training. To mitigate issues such as kernel death during the loading phase, we utilized checkpoints.

2. **S3:**

We stored our image data in an S3 bucket along with the annotations. The bucket was accessed by the EC2 instances for training purposes and ultimately for the validation. Proper permissions and roles (LabRole) were assigned to the EC2 instances to enable seamless access to the S3 storage.

3. **AWS Lambda:**

This service played a crucial role in deploying and utilizing the trained model. We tried to configure Lambda to trigger the prediction code at specified intervals. To achieve this, we set appropriate roles to access the S3 bucket (LabRole) and configured a timeout of 1 hour for the benchmark computation.

4. **DynamoDB:**

We explored using DynamoDB to record and store the benchmarks obtained during the training and testing phases in a table that we created called *"benchmark"*. This allowed us to track and analyze the performance of our pipeline effectively.

5. Development

For the development of our cloud application we proceeded in many steps with different turnbacks and errors encountered, here we will talk about the process for each of the services used

- **EC2 Instance:**

After many trials with different kind of machines, we opted to set up an EC2 instance with the following specifications:

- Instance Type: **t2.large**
- RAM: 8 GB
- vCPUs: 2

In this way we managed to surpass many problems that we had in installing the dependencies given the low RAM of the other instances given that we tried to be cost efficient at most.

Below the details of our chosen instance:

The screenshot displays the AWS Management Console interface for an EC2 instance. At the top, a table lists instances, with the selected instance 'i-06f83e57707a88396' highlighted. Below this, the 'Dettagli' (Details) tab is active, showing various configuration parameters for the instance 'i-06f83e57707a88396 (large)'. The details are organized into three columns:

- Left Column (General Information):**
 - ID istanza: i-06f83e57707a88396
 - Indirizzo IPv6: --
 - Tipo di nome host: Nome IP: ip-172-31-42-29.ec2.internal
 - Indirizzo IP assegnato automaticamente: 44.201.84.103 (IP pubblico)
 - Ruolo IAM: LabRole
 - Indirizzo IP privato: --
 - Indirizzo IP pubblico: 44.201.84.103
 - Indirizzo IP privato: 172.31.82.29
 - Indirizzo IP pubblico: ec2-44-201-84-103.compute-1.amazonaws.com
 - Indirizzo IP elastici: --
 - Risultato di AWS Compute Optimizer: Consenso esplicito ad AWS Compute Optimizer per le raccomandazioni
 - Nome del gruppo di Auto Scaling: --
- Middle Column (Network and Security):**
 - Indirizzo IPv4 pubblico: 44.201.84.103
 - Indirizzo IPv4 privato: 172.31.82.29
 - Indirizzo IPv6 pubblico: ec2-44-201-84-103.compute-1.amazonaws.com
 - Indirizzo IPv6 privato: --
 - Indirizzo IPv4 elastici: --
 - Indirizzo IPv6 elastici: --
 - Indirizzo IPv4 pubblico: ec2-44-201-84-103.compute-1.amazonaws.com
 - Indirizzo IPv6 pubblico: --
 - Indirizzo IPv4 elastici: --
 - Indirizzo IPv6 elastici: --
- Right Column (Advanced Settings):**
 - Indirizzo IPv4 pubblico: 44.201.84.103
 - Indirizzo IPv4 privato: 172.31.82.29
 - Indirizzo IPv6 pubblico: ec2-44-201-84-103.compute-1.amazonaws.com
 - Indirizzo IPv6 privato: --
 - Indirizzo IPv4 elastici: --
 - Indirizzo IPv6 elastici: --
 - Indirizzo IPv4 pubblico: ec2-44-201-84-103.compute-1.amazonaws.com
 - Indirizzo IPv6 pubblico: --
 - Indirizzo IPv4 elastici: --
 - Indirizzo IPv6 elastici: --

- **Environment:**

To train the model that is used as the final validator we had to create an empty **conda environment** (*py_env*) and install on it the dependencies needed:

- **Torch 2.0.1**

The most important package that we needed to install in order to train, test and then validate the model, the whole concept of the project revolves around this.

- Pandas 2.0.3

- Numpy 1.25.21

- Seaborn

- Matplotlib

- **Boto3**

Which allowed us to communicate both with the S3 bucket and the DynamoDB.

- Pickle

We installed **Jupyter Notebook** thanks to *conda* to have a better user experience in the development of the code to train our model.

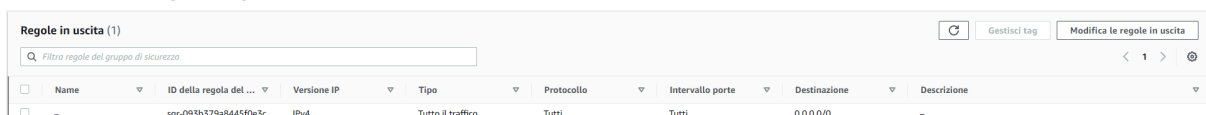
To connect the EC2 to an external URL we created a new security group with different rules to enable ingoing...



The screenshot shows the 'Regole in entrata (3)' (Inbound Rules) section of the AWS IAM console. It displays a table with three rules. Each rule has a checkbox, a name, an ID, a version, a type, a protocol, a port range, an origin, and a description. The rules are for TCP, HTTPS, and SSH.

<input type="checkbox"/>	Name	ID della regola del ...	Versione IP	Tipo	Protocollo	Intervallo porte	Origine	Descrizione
<input type="checkbox"/>	-	sgr-061567f6e11d278a3	IPv4	Regola TCP personaliz...	TCP	8888	0.0.0.0/0	-
<input type="checkbox"/>	-	sgr-00a65321e7bcc5b75	IPv4	HTTPS	TCP	443	0.0.0.0/0	-
<input type="checkbox"/>	-	sgr-01178c6507afcfcc3	IPv4	SSH	TCP	22	0.0.0.0/0	-

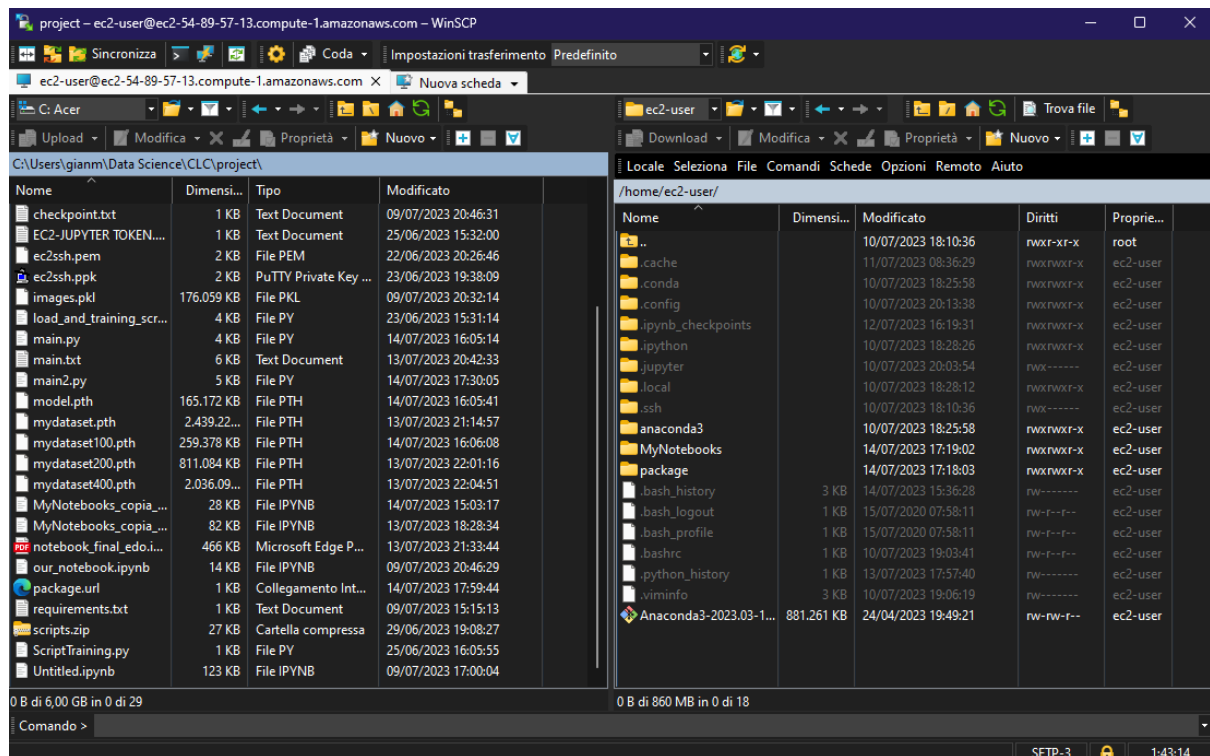
...and outgoing traffic



The screenshot shows the 'Regole in uscita (1)' (Outbound Rules) section of the AWS IAM console. It displays a table with one rule. The rule has a checkbox, a name, an ID, a version, a type, a protocol, a port range, a destination, and a description. The rule is for 'Tutto il traffico' (All traffic).

<input type="checkbox"/>	Name	ID della regola del ...	Versione IP	Tipo	Protocollo	Intervallo porte	Destinazione	Descrizione
<input type="checkbox"/>	-	sgr-093b379a8445f0e3c	IPv4	Tutto il traffico	Tutti	Tutti	0.0.0.0/0	-

To make it easy to manage files on the instance we used a third party software called **WinSCP**:



To connect we used the keys stored in the “.pem” file extracted upon the creation of our EC2.

- **S3 Bucket:**

The creation of the S3 Bucket was very smooth, in this we uploaded the images and the annotations as well as the model and the dataset (image preprocessed).

To access the S3 bucket from our newly created instance we had to grant the appropriate permissions (**LabInstanceProfile**) to the instance and modify its policy to allow access to the files.

Now that our environment is ready and the services are well aligned we can proceed to the pre-processing of the data and the training of our model.

- **Training**

We made the decision to switch from using the YOLOv5 architecture to the ResNet34 architecture as our pre-trained model.

The primary reason for this change was the difficulties we encountered during the installation of the dependencies required for YOLOv5. Despite following the installation guidelines, the YOLOv5 package was unable to locate the installed dependencies.

We suspect that this issue stemmed from a conflict between library different versions.

To proceed with our project and ensure a viable alternative, we opted to explore the ResNet34 architecture. ResNet34 has been widely used in the literature for similar tasks and has proven to be effective in various image classification tasks, including mask-wearing behavior classification.

By switching to ResNet34, we aimed to leverage the pre-trained model's capabilities and benefit from its established performance. This decision was driven by the need to overcome the challenges faced with YOLOv5 and continue our project smoothly.

Thus the problem is not of object detection but of object classification, where in each image with multiple masks, multiple objects to classify are derived.

Once we obtain the dataset with the tuples of tensors representing the object to identify, we proceed to train our model as follows:

```
1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4
5 model=models.resnet34(pretrained=True)
6 criterion=torch.nn.CrossEntropyLoss()
7 optimizer=optim.SGD(model.parameters(),lr=0.001,momentum=0.9)
8
9 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
10
11 def train(model, criterion, optimizer, train_loader, val_loader, num_epochs):
12     model = model.to(device)
13     criterion = criterion.to(device)
14
15     for epoch in range(num_epochs):
16         model.train()
17         train_loss = 0.0
18         train_correct = 0
19         train_total = 0
20
21         for inputs, labels in train_loader:
22             inputs = inputs.to(device)
23             labels = labels.to(device)
24
25             optimizer.zero_grad()
26
27             # forward pass
28             outputs = model(inputs)
29             loss = criterion(outputs, labels)
30
31             # backward pass and optimization
32             loss.backward()
33             optimizer.step()
34
35             train_loss += loss.item()
36             _, predicted = torch.max(outputs.data, 1)
37             train_total += labels.size(0)
38             train_correct += (predicted == labels).sum().item()
39
40         train_loss /= len(train_loader)
41         train_accuracy = 100.0 * train_correct / train_total
42
43         # validation
44         model.eval()
45         val_loss = 0.0
46         val_correct = 0
47         val_total = 0
48
49         with torch.no_grad():
50             for inputs, labels in val_loader:
51                 inputs = inputs.to(device)
52                 labels = labels.to(device)
53
54                 outputs = model(inputs)
55                 loss = criterion(outputs, labels)
56
57                 val_loss += loss.item()
58                 _, predicted = torch.max(outputs.data, 1)
59                 val_total += labels.size(0)
60                 val_correct += (predicted == labels).sum().item()
61
62         val_loss /= len(val_loader)
63         val_accuracy = 100.0 * val_correct / val_total
64
65         # print epoch statistics
66         print(f"Epoch {epoch+1}/{num_epochs}:")
67         print(f"Train Loss: {train_loss:.4f} | Train Accuracy: {train_accuracy:.2f}%")
68         print(f"Val Loss: {val_loss:.4f} | Val Accuracy: {val_accuracy:.2f}%")
69         print("-----")
70
71     print("Training complete.")
72
73 num_epochs = 5
74 learning_rate = 0.0001
75 optimizer = optim.SGD(model.parameters(),lr=learning_rate,momentum=0.9)
76
77 start_time = time.time()
78 train(model, criterion, optimizer, train_loader, val_loader, num_epochs)
79 elapsed_time = time.time() - start_time
```

Given that the pre-trained was already giving high accuracy in the first epoch, we chose to train it for 5 epochs overall, obtaining very satisfying results given a train set of 60%, test and val set both of 20%:

Throughout this phase of training, we encountered challenges related to resource limitations, which necessitated scaling the EC2 instance. We also implemented checkpoints to ensure progress could be saved in case of kernel death.

We trained the model using the concept of transfer knowledge, so we took the pretrained model, we froze the first 6 layers and finally we fine tuned the model.

```
Epoch 1/5:
Train Loss: 1.2388 | Train Accuracy: 82.56%
Val Loss: 0.2259 | Val Accuracy: 95.70%
-----
Epoch 2/5:
Train Loss: 0.0874 | Train Accuracy: 96.89%
Val Loss: 0.1521 | Val Accuracy: 95.70%
-----
Epoch 3/5:
Train Loss: 0.0422 | Train Accuracy: 98.77%
Val Loss: 0.1379 | Val Accuracy: 95.58%
-----
Epoch 4/5:
Train Loss: 0.0151 | Train Accuracy: 99.80%
Val Loss: 0.1406 | Val Accuracy: 95.95%
-----
Epoch 5/5:
Train Loss: 0.0144 | Train Accuracy: 99.71%
Val Loss: 0.1480 | Val Accuracy: 95.70%
-----
Training complete.
```

Thanks to torch utilities we saved our model in a file called “*model.pth*” to keep track of it and to easily pass it and be used by different sources (AWS Lambda).

- **AWS Lambda:**

The main idea in our project was to use AWS Lambda to test our model by triggering the activation thanks to an **EventBridge** and to perform an identification for a given set of images and to store some values inside a **DynamoDB** directly from the AWS Lambda function.

Unfortunately we encountered many problems in this part of the development and couldn't resolve them all, but we still want to illustrate the main idea behind this piece of our architecture and the work done.

Given that the code we wanted to be executed on the AWS Lambda needed dependencies such as torch, we had to deploy a package with a 'ready-to-go' environment (*landa_env*) and a main file with the testing in clear.

To do so we created a folder called 'package' with the structure:

```
| code
  | main.py
  | unzip_torch.py
| lambda_env
| deploy.sh
```

In the code repository we can find the **main.py** which contains the famous **lambda_handler**, i.e. the function that will be executed when the EventBridge is triggered:

```

47 def main():
48     # Load model from s3 bucket
49     s3_client.download_file(bucket_name, file_key, 'model.pth')
50     with open('model.pth', 'rb') as file:
51         model = torch.load(file)
52
53
54     # Load chunks of dataset
55     s3_client.download_file(bucket_name, 'mydataset100.pth', 'mydataset100.pth')
56     with open('mydataset100.pth', 'rb') as file:
57         mydataset100 = torch.load(file)
58
59     s3_client.download_file(bucket_name, 'mydataset200.pth', 'mydataset200.pth')
60     with open('mydataset200.pth', 'rb') as file:
61         mydataset200 = torch.load(file)
62
63     # s3_client.download_file(bucket_name, 'mydataset300.pth', 'mydataset300.pth')
64     # with open('mydataset300.pth', 'rb') as file:
65     #     mydataset300 = torch.load(file)
66
67     # s3_client.download_file(bucket_name, 'mydataset400.pth', 'mydataset400.pth')
68     # with open('mydataset400.pth', 'rb') as file:
69     #     mydataset400 = torch.load(file)
70
71
72
73     data_loader_100 = DataLoader(dataset=mydataset100, batch_size=32, shuffle=True, num_workers=1)
74
75     data_loader_200 = DataLoader(dataset=mydataset200, batch_size=32, shuffle=True, num_workers=1)
76
77     data_loader_400 = DataLoader(dataset=mydataset400, batch_size=32, shuffle=True, num_workers=1)
78
79     model.eval() # Set model to evaluation mode
80     criterion=torch.nn.CrossEntropyLoss()
81
82
83     # 100
84     val_loss = 0.0
85     val_correct = 0
86     val_total = 0
87
88     start_time = time.time()
89
90     with torch.no_grad():
91         for images, labels in data_loader_100:
92             outputs = model(images)
93
94             loss = criterion(outputs, labels)
95             # Compute validation accuracy
96             _, predicted = torch.max(outputs.data, 1)
97             val_total += labels.size(0)
98             val_correct += (predicted == labels).sum().item()
99             val_loss_100 += loss.item() * images.size(0)
100
101     val_loss_100 /= len(data_loader_100.dataset)
102     val_accuracy_100 = 100.0 * val_correct/val_total
103     elapsed_time_100 = time.time() - start_time
104
105     item100 = {
106         'size': len(data_loader_100.dataset),
107         'true_prediction': val_accuracy_100,
108         'time': elapsed_time_100
109     }
110     table.put_item(Item=item100)
111
112     # 200
113     val_loss = 0.0
114     val_correct = 0
115     val_total = 0
116
117     start_time = time.time()
118
119     with torch.no_grad():
120         for images, labels in data_loader_200:
121             outputs = model(images)
122
123             loss = criterion(outputs, labels)
124             # Compute validation accuracy
125             _, predicted = torch.max(outputs.data, 1)
126             val_total += labels.size(0)
127             val_correct += (predicted == labels).sum().item()
128             val_loss_200 += loss.item() * images.size(0)
129
130     val_loss_200 /= len(data_loader_200.dataset)
131     val_accuracy_200 = 100.0 * val_correct/val_total
132     elapsed_time_200 = time.time() - start_time
133
134     item200 = {
135         'size': len(data_loader_200.dataset),
136         'true_prediction': val_accuracy_200,
137         'time': elapsed_time_200
138     }
139     table.put_item(Item=item200)
140

```

```

140
141 # 400
142 val_loss = 0.0
143 val_correct = 0
144 val_total = 0
145
146 start_time = time.time()
147
148 with torch.no_grad():
149     for images, labels in data_loader_400:
150         outputs = model(images)
151
152         loss = criterion(outputs, labels)
153         # Compute validation accuracy
154         _, predicted = torch.max(outputs.data, 1)
155         val_total += labels.size(0)
156         val_correct += (predicted == labels).sum().item()
157         val_loss_400 += loss.item() * images.size(0)
158
159 val_loss_400 /= len(data_loader_400.dataset)
160 val_accuracy_400 = 100.0 * val_correct / val_total
161 elapsed_time_400 = time.time() - start_time
162
163 item400 = {
164     'size': len(data_loader_400.dataset),
165     'true_prediction': val_accuracy_400,
166     'time': elapsed_time_400
167 }
168 table.put_item(Item=item400)
169
170
171 def lambda_handler(event, context):
172     if __name__ == '__main__':
173         main()

```

The **unzip_torch.py** is a file thinked to fit the package in the size limit (250MB) of deployment package, the idea was to compress the torch library and at runtime decompress it under a '/tmp' directory which grants an additional 500MB of storage.

Finally, the **deploy.sh** is a script built to (of course) deploy the package to the AWS Lambda function.

The main problem we found here was the size of the package, although the management of the memory of torch's dependencies we couldn't diminish the size of our package even after compressing it to a zip.

And when we tried to download it or to put it on the S3 bucket to be accessible the AWS couldn't respond in any way resulting in the interruption of the connection.

- **Alternative solution:**

Instead of using the AWS Lambda, we used the same code (except for the lambda_handler function of course) in a python file inside the EC2 and through the notebook we managed to call the testing for the given set of images.

The set of images we chose to make the test on are three of the size, respectively, of 100, 200 and 400 images.

Each image could contain more masks than just one, so the detection can be performed more than once per image to detect multiple masks.

For every set we want to perform a test and gather information about it using the model we talked about earlier.

The properties in which we are interested are:

- **Size** of the set (represented in number of objects to detect, not in images)
- **Time** took by the model to analyze the whole set
- Percentage of **predictions** made by the model, indicating its efficiency

These will be stored in different records inside the **DynamoDB** table that we early created, called *benchmark*, where the partition key is the size of the set.

To access the DynamoDB was sufficient to instantiate a client thanks to boto3 and then upload the records using the JSON format:

```
124 response = dynamodb_client.put_item(  
125     TableName=table_name,  
126     Item = {  
127         'size': {"S":str(len(data_loader_400.dataset))},  
128         'true_prediction': {"S":str(val_accuracy_400)},  
129         'time' : {"S":str(elapsed_time_400)}  
130     },  
131 )
```

The results obtained can be seen in the table below, extracted directly from our DynamoDB “benchmark”.

size	time	true_prediction
433	66.63697910308838	98.84526558891454
3399	463.7631437778473	98.29361576934393
1354	195.2414755821228	98.59675036927622

As we can see the results obtained are in line with what we could expect especially for the prediction.

6. Conclusion

Taking a look at the results obtained it seems that they are in line with what we expected, the increment in time of execution scales linearly with the size of each set of analysis.

Moreover the percentage of correct prediction is very high and remains constant even if the samples to analyze are much bigger, implying a good reliability of the model and the project in general.

The cost efficiency has also been a key aspect of our development and we managed to always stay at low costs for the requested services, also thanks to the free tier plan that allowed us to use many resources for free.

Some limitations of the Learner Lab have been a bit harsh to handle, such as the impossibility of creating an IAM role that we wanted to use to make the instances and services available for different users.

But these kind of obstacles have been easily overridden and we can affirm that the Learner Lab is an excellent tool to take confidence with the most used cloud infrastructures.

In conclusion we can say that this project provided for us a valuable opportunity to acquire and apply essential tools in the realm of cloud technologies, which have a significant impact on our professional lives.

Our choice to undertake this challenging project aimed at gaining a comprehensive understanding of potential obstacles and improving our skills. Even the challenges we faced during the Lambda AWS segment served as valuable learning experiences, helping us grasp the associated difficulties and identify areas for further improvement and learning.